

Detection of Fraudulent Transactions

Ayoub HAIDA

April 3, 2023

1 Detection of Fraudulent Transactions

The objective of this case study is to employ different models based on classification techniques for the purpose of identifying whether a given transaction is a regular payment or a fraudulent one.

2 Problem Definition

In this case, the response variable takes a value of 1 in case the given transaction is fraud and 0 otherwise.

The dataset provided contains information about credit card transactions carried out by European cardholders in September 2013. The data pertains to transactions that took place over a period of two days, during which there were a total of 492 fraudulent transactions out of 284,807. It is important to note that this dataset is highly imbalanced, with only 0.172% of all transactions being identified as frauds. The main objective is to predict instances of fraud. The response variable, 'Class', takes on a value of 1 if the transaction is fraudulent and 0 otherwise. The features in the dataset are a result of PCA transformation and may not be immediately interpretable based on their names.

Here, you can find the dataset details as well as download it: <https://www.kaggle.com/mlg-ulb/creditcardfraud>

3 Loading the Packages and the Data

```
[1]: import numpy as np
import pandas as pd
from matplotlib import pyplot
from pandas import read_csv, set_option
from pandas.plotting import scatter_matrix
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, KFold, cross_val_score, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
```

```

from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.pipeline import Pipeline
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier, \
    RandomForestClassifier, ExtraTreesClassifier
from sklearn.metrics import classification_report, confusion_matrix, \
    accuracy_score
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasClassifier
from keras.optimizers import SGD
from pickle import dump
from pickle import load

```

```

[2]: # load dataset
dataset = read_csv('/kaggle/input/creditcardfraud/creditcard.csv')

```

```

[3]: import warnings
warnings.filterwarnings('ignore')

```

4 Exploratory Data Analysis

4.1 Stats

```

[4]: # shape
dataset.shape

```

```

[4]: (284807, 31)

```

```

[5]: set_option('display.width', 100)
dataset.head()

```

```

[5]:   Time      V1      V2      V3      V4      V5      V6      V7
V8      V9 \
0   0.0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599
0.098698  0.363787
1   0.0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803
0.085102 -0.255425
2   1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461
0.247676 -1.514654
3   1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609
0.377436 -1.387024
4   2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941
-0.270533  0.817739

...      V21      V22      V23      V24      V25      V26      V27
V28 Amount \

```

```

0 ... -0.018307  0.277838 -0.110474  0.066928  0.128539 -0.189115  0.133558
-0.021053  149.62
1 ... -0.225775 -0.638672  0.101288 -0.339846  0.167170  0.125895 -0.008983
0.014724    2.69
2 ...  0.247998  0.771679  0.909412 -0.689281 -0.327642 -0.139097 -0.055353
-0.059752  378.66
3 ... -0.108300  0.005274 -0.190321 -1.175575  0.647376 -0.221929  0.062723
0.061458  123.50
4 ... -0.009431  0.798278 -0.137458  0.141267 -0.206010  0.502292  0.219422
0.215153   69.99

```

```

      Class
0         0
1         0
2         0
3         0
4         0

```

[5 rows x 31 columns]

```
[6]: set_option('display.max_rows', 500)
      dataset.dtypes
```

```

[6]: Time      float64
     V1        float64
     V2        float64
     V3        float64
     V4        float64
     V5        float64
     V6        float64
     V7        float64
     V8        float64
     V9        float64
     V10       float64
     V11       float64
     V12       float64
     V13       float64
     V14       float64
     V15       float64
     V16       float64
     V17       float64
     V18       float64
     V19       float64
     V20       float64
     V21       float64
     V22       float64
     V23       float64

```

```

V24      float64
V25      float64
V26      float64
V27      float64
V28      float64
Amount   float64
Class    int64
dtype: object

```

```

[7]: pd.set_option('display.float_format', '{:.3f}'.format)
      dataset.describe()

```

```

[7]:
      Time      V1      V2      V3      V4      V5
V6      V7 \
count 284807.000 284807.000 284807.000 284807.000 284807.000 284807.000
284807.000 284807.000
mean   94813.860    0.000    0.000   -0.000    0.000    0.000
0.000   -0.000
std    47488.146    1.959    1.651    1.516    1.416    1.380
1.332    1.237
min      0.000   -56.408   -72.716   -48.326   -5.683   -113.743
-26.161   -43.557
25%    54201.500   -0.920   -0.599   -0.890   -0.849   -0.692
-0.768   -0.554
50%    84692.000    0.018    0.065    0.180   -0.020   -0.054
-0.274    0.040
75%   139320.500    1.316    0.804    1.027    0.743    0.612
0.399    0.570
max   172792.000    2.455    22.058    9.383   16.875   34.802
73.302   120.589

      V8      V9  ...      V21      V22      V23      V24
V25 \
count 284807.000 284807.000  ... 284807.000 284807.000 284807.000 284807.000
284807.000
mean      0.000   -0.000  ...    0.000   -0.000    0.000    0.000
0.000
std      1.194    1.099  ...    0.735    0.726    0.624    0.606
0.521
min     -73.217  -13.434  ...   -34.830  -10.933  -44.808   -2.837
-10.295
25%     -0.209   -0.643  ...   -0.228   -0.542   -0.162   -0.355
-0.317
50%      0.022   -0.051  ...   -0.029    0.007   -0.011    0.041
0.017
75%      0.327    0.597  ...    0.186    0.529    0.148    0.440
0.351

```

max	20.007	15.595	...	27.203	10.503	22.528	4.585
7.520							

	V26	V27	V28	Amount	Class
count	284807.000	284807.000	284807.000	284807.000	284807.000
mean	0.000	-0.000	-0.000	88.350	0.002
std	0.482	0.404	0.330	250.120	0.042
min	-2.605	-22.566	-15.430	0.000	0.000
25%	-0.327	-0.071	-0.053	5.600	0.000
50%	-0.052	0.001	0.011	22.000	0.000
75%	0.241	0.091	0.078	77.165	0.000
max	3.517	31.612	33.848	25691.160	1.000

[8 rows x 31 columns]

To begin with, we will examine the number of instances of fraud as compared to non-fraud cases in the dataset.


```
[8]: class_names = {0: 'Not Fraud', 1: 'Fraud'}
      print(dataset.Class.value_counts().rename(index = class_names))
```

```
Not Fraud    284315
Fraud         492
Name: Class, dtype: int64
```

The distribution of instances in the dataset is skewed, with a significant majority of the transactions being classified as non-fraudulent.

4.2 Visualization

```
[9]: # histograms
      dataset.hist(sharex=False, sharey=False, xlabelsize=1, ylabelsize=1,
                  figsize=(12,12))
      pyplot.show()
```



4.3 Data Preparation

```
[10]: print('Null Values =',dataset.isnull().values.any())
```

Null Values = False

There is no null in the data.

```
[11]: from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2

bestfeatures = SelectKBest( k=10)
bestfeatures
```

```

Y= dataset["Class"]
X = dataset.loc[:, dataset.columns != 'Class']
fit = bestfeatures.fit(X,Y)
dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(X.columns)

featureScores = pd.concat([dfcolumns,dfscores],axis=1)
featureScores.columns = ['Specs','Score']
print(featureScores.nlargest(10,'Score'))  #print 10 best features

```

	Specs	Score
17	V17	33979.169
14	V14	28695.548
12	V12	20749.822
10	V10	14057.980
16	V16	11443.349
3	V3	11014.508
7	V7	10349.605
11	V11	6999.355
4	V4	5163.832
18	V18	3584.381

While certain features in the dataset may be important for detecting fraud, the process of selecting these features has not been emphasized or prioritized.

5 Algorithms and Models

5.1 Train-Test-Split and Evaluation Metrics

```

[14]: Y= dataset["Class"]
X = dataset.loc[:, dataset.columns != 'Class']
validation_size = 0.2
seed = 7
X_train, X_validation, Y_train, Y_validation = train_test_split(X, Y,
    ↪test_size=validation_size, random_state=seed)
scoring = 'accuracy'

```

```

[16]: num_folds = 10
seed = 7

```

```

[17]: models = []
models.append(('LR', LogisticRegression()))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('SVM', SVC()))
models.append(('NN', MLPClassifier()))


```

```
models.append(('AB', AdaBoostClassifier()))
models.append(('GBM', GradientBoostingClassifier()))
models.append(('RF', RandomForestClassifier()))
models.append(('ET', ExtraTreesClassifier()))
```

```
[18]: results = []
names = []
for name, model in models:
    kfold = KFold(n_splits=num_folds)
    cv_results = cross_val_score(model, X_train, Y_train, cv=kfold,
    ↪scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
```

```
LR: 0.998885 (0.000205)
LDA: 0.999364 (0.000136)
KNN: 0.998310 (0.000290)
CART: 0.999109 (0.000245)
NB: 0.992745 (0.000733)
SVM: 0.998280 (0.000278)
NN: 0.998429 (0.000361)
AB: 0.999192 (0.000154)
GBM: 0.999008 (0.000285)
RF: 0.999570 (0.000107)
ET: 0.999565 (0.000122)
```

```
[20]: # compare algorithms
fig = pyplot.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
pyplot.boxplot(results)
ax.set_xticklabels(names)
fig.set_size_inches(15,7)
pyplot.show()
```

```
[21]: model = DecisionTreeClassifier()
      model.fit(X_train, Y_train)
```

```
[21]: DecisionTreeClassifier()
```

```
[22]: rescaledValidationX = X_validation
      predictions = model.predict(rescaledValidationX)
      print(accuracy_score(Y_validation, predictions))
      print(confusion_matrix(Y_validation, predictions))
      print(classification_report(Y_validation, predictions))
```

```
0.9991397773954567
```


```
[[56833  29]
```

```
 [  20  80]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56862
1	0.73	0.80	0.77	100
accuracy			1.00	56962
macro avg	0.87	0.90	0.88	56962
weighted avg	1.00	1.00	1.00	56962

```
[23]: df_cm = pd.DataFrame(confusion_matrix(Y_validation, predictions), columns=np.
      ↪unique(Y_validation), index = np.unique(Y_validation))
      df_cm.index.name = 'Actual'
      df_cm.columns.name = 'Predicted'
```

```
sns.heatmap(df_cm, cmap="Blues", annot=True,annot_kws={"size": 16});
```



Despite achieving good results, it is worth noting that 20 out of 100 fraud cases are not detected. Therefore, it is crucial to prioritize the metric of recall, which minimizes false negatives.

5.2 Model Tuning

```
[24]: scoring = 'recall'
```

```
[25]: models = []
models.append(('LR', LogisticRegression()))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('SVM', SVC()))
models.append(('NN', MLPClassifier()))
models.append(('AB', AdaBoostClassifier()))
models.append(('GBM', GradientBoostingClassifier()))
models.append(('RF', RandomForestClassifier()))
models.append(('ET', ExtraTreesClassifier()))
```

```
[ ]: results = []
names = []
for name, model in models:
    kfold = KFold(n_splits=num_folds)
    cv_results = cross_val_score(model, X_train, Y_train, cv=kfold,
    ↪scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
```

```
LR: 0.631287 (0.116650)
LDA: 0.758283 (0.045450)
KNN: 0.023882 (0.019671)
CART: 0.749286 (0.070769)
NB: 0.659465 (0.080686)
SVM: 0.000000 (0.000000)
```

Given the LDA has the best recall out of all the models, it is used to evaluate the test set

```
[ ]: model = LinearDiscriminantAnalysis()
model.fit(X_train, Y_train)
```

```
[185]: rescaledValidationX = X_validation
predictions = model.predict(rescaledValidationX)
print(accuracy_score(Y_validation, predictions))
print(confusion_matrix(Y_validation, predictions))
print(classification_report(Y_validation, predictions))
```

```
0.9995435553526912
```

```
[[56854    8]
```


```
[   18   82]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56862
1	0.91	0.82	0.86	100
accuracy			1.00	56962
macro avg	0.96	0.91	0.93	56962

weighted avg 1.00 1.00 1.00 56962

```
[186]: df_cm = pd.DataFrame(confusion_matrix(Y_validation, predictions), columns=np.  
    ↪unique(Y_validation), index = np.unique(Y_validation))  
df_cm.index.name = 'Actual'  
df_cm.columns.name = 'Predicted'  
sns.heatmap(df_cm, cmap="Blues", annot=True,annot_kws={"size": 16});
```

```
[186]: <matplotlib.axes._subplots.AxesSubplot at 0x20b99399978>
```



```
[204]: Y_train.head()
```

```
[204]: 44828      0  
221877      0  
278826      0  
149792      0  
226041      0  
Name: Class, dtype: int64
```

```
[41]: df = pd.concat([X_train, Y_train], axis=1)  
  
fraud_df = df.loc[df['Class'] == 1]
```

```

non_fraud_df = df.loc[df['Class'] == 0][:492]# fraud classes = 492.

normal_distributed_df = pd.concat([fraud_df, non_fraud_df])

df_new = normal_distributed_df.sample(frac=1, random_state=42)
Y_train_new= df_new["Class"]
X_train_new = df_new.loc[:, dataset.columns != 'Class']

dataset.head()

```

```

[41]:   Time      V1      V2      V3      V4      V5      V6      V7      V8      V9  ...
      V21      V22      V23  \
0   0.0 -1.360 -0.073  2.536  1.378 -0.338  0.462  0.240  0.099  0.364  ...
      -0.018  0.278 -0.110
1   0.0  1.192  0.266  0.166  0.448  0.060 -0.082 -0.079  0.085 -0.255  ...
      -0.226 -0.639  0.101
2   1.0 -1.358 -1.340  1.773  0.380 -0.503  1.800  0.791  0.248 -1.515  ...
      0.248  0.772  0.909
3   1.0 -0.966 -0.185  1.793 -0.863 -0.010  1.247  0.238  0.377 -1.387  ...
      -0.108  0.005 -0.190
4   2.0 -1.158  0.878  1.549  0.403 -0.407  0.096  0.593 -0.271  0.818  ...
      -0.009  0.798 -0.137

      V24      V25      V26      V27      V28  Amount  Class
0  0.067  0.129 -0.189  0.134 -0.021  149.62      0
1 -0.340  0.167  0.126 -0.009  0.015    2.69      0
2 -0.689 -0.328 -0.139 -0.055 -0.060  378.66      0
3 -1.176  0.647 -0.222  0.063  0.061  123.50      0
4  0.141 -0.206  0.502  0.219  0.215   69.99      0

[5 rows x 31 columns]

```

```

[42]: print('Distribution of the Classes in the subsample dataset')
      print(df_new['Class'].value_counts()/len(df_new))
      sns.countplot('Class', data=df_new)
      pyplot.title('Equally Distributed Classes', fontsize=14)
      pyplot.show()

```


Distribution of the Classes in the subsample dataset

```

1    0.5
0    0.5

```

Name: Class, dtype: float64



```
[43]: scoring='accuracy'
```

```
[44]: # spot check the algorithms
models = []
models.append(('LR', LogisticRegression()))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('SVM', SVC()))
models.append(('NN', MLPClassifier()))
models.append(('AB', AdaBoostClassifier()))
models.append(('GBM', GradientBoostingClassifier()))
models.append(('RF', RandomForestClassifier()))
models.append(('ET', ExtraTreesClassifier()))
```

```
[45]: EnableDLModelsFlag = 1
if EnableDLModelsFlag == 1 :
    def create_model(neurons=12, activation='relu', learn_rate = 0.01,
        ↪momentum=0):
        # create model
        model = Sequential()
```

```

        model.add(Dense(X_train.shape[1], input_dim=X_train.shape[1],
↪activation=activation))
        model.add(Dense(32, activation=activation))
        model.add(Dense(1, activation='sigmoid'))
        optimizer = SGD(lr=learn_rate, momentum=momentum)
        model.compile(loss='binary_crossentropy', optimizer='adam',
↪metrics=['accuracy'])
        return model
    models.append(('DNN', KerasClassifier(build_fn=create_model, epochs=50,
↪batch_size=10, verbose=0)))

```

```

[46]: results = []
names = []
for name, model in models:
    kfold = KFold(n_splits=num_folds, random_state=seed)
    cv_results = cross_val_score(model, X_train_new, Y_train_new, cv=kfold,
↪scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

```

LR: 0.931911 (0.024992)

LDA: 0.905473 (0.027422)

KNN: 0.648258 (0.044550)

CART: 0.907565 (0.022669)

NB: 0.860771 (0.027234)

SVM: 0.522356 (0.048395)

NN: 0.648712 (0.100137)

AB: 0.924830 (0.024068)

GBM: 0.934982 (0.015132)

RF: 0.932931 (0.015859)

ET: 0.931962 (0.031043)

WARNING:tensorflow:From D:\Anaconda\lib\site-packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future

version.

Instructions for updating:

Colocations handled automatically by placer.


WARNING:tensorflow:From D:\Anaconda\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

DNN: 0.498011 (0.050742)

```
[47]: fig = pyplot.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
pyplot.boxplot(results)
ax.set_xticklabels(names)
fig.set_size_inches(8,4)
pyplot.show()
```



Since GBM has been determined as the top-performing model among all other models, a grid search is conducted to optimize the GBM model's performance by varying the number of estimators and

maximum depth parameters.

```
[48]: n_estimators = [20,180,1000]
max_depth= [2, 3,5]
param_grid = dict(n_estimators=n_estimators, max_depth=max_depth)
model = GradientBoostingClassifier()
kfold = KFold(n_splits=num_folds, random_state=seed)
grid = GridSearchCV(estimator=model, param_grid=param_grid, scoring=scoring,
                    cv=kfold)
grid_result = grid.fit(X_train_new, Y_train_new)

#Print Results
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
ranks = grid_result.cv_results_['rank_test_score']
for mean, stdev, param, rank in zip(means, stds, params, ranks):
    print("#%d %f (%f) with: %r" % (rank, mean, stdev, param))
```

Best: 0.936992 using {'max_depth': 5, 'n_estimators': 1000}

#3 0.931911 (0.016958) with: {'max_depth': 2, 'n_estimators': 20}

#6 0.929878 (0.017637) with: {'max_depth': 2, 'n_estimators': 180}

#9 0.924797 (0.021358) with: {'max_depth': 2, 'n_estimators': 1000}

#6 0.929878 (0.020476) with: {'max_depth': 3, 'n_estimators': 20}

#3 0.931911 (0.011120) with: {'max_depth': 3, 'n_estimators': 180}

#3 0.931911 (0.017026) with: {'max_depth': 3, 'n_estimators': 1000}

#8 0.928862 (0.022586) with: {'max_depth': 5, 'n_estimators': 20}

#2 0.934959 (0.015209) with: {'max_depth': 5, 'n_estimators': 180}

#1 0.936992 (0.012639) with: {'max_depth': 5, 'n_estimators': 1000}

```
[49]: model = GradientBoostingClassifier(max_depth= 5, n_estimators = 1000)
model.fit(X_train_new, Y_train_new)
```

```
[49]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
                                learning_rate=0.1, loss='deviance', max_depth=5,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
```

```

min_weight_fraction_leaf=0.0, n_estimators=1000,
n_iter_no_change=None, presort='auto',
random_state=None, subsample=1.0, tol=0.0001,
validation_fraction=0.1, verbose=0,
warm_start=False)

```

```

[50]: predictions = model.predict(X_validation)
print(accuracy_score(Y_validation, predictions))
print(confusion_matrix(Y_validation, predictions))
print(classification_report(Y_validation, predictions))

```

0.9668199852533268

```
[[54972  1890]
```

```
[   0   100]]
```


	precision	recall	f1-score	support
0	1.00	0.97	0.98	56862
1	0.05	1.00	0.10	100
accuracy			0.97	56962
macro avg	0.53	0.98	0.54	56962
weighted avg	1.00	0.97	0.98	56962

```

[51]: df_cm = pd.DataFrame(confusion_matrix(Y_validation, predictions), columns=np.
    ↪unique(Y_validation), index = np.unique(Y_validation))
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'
sns.heatmap(df_cm, cmap="Blues", annot=True,annot_kws={"size": 16});

```

```
[51]: <matplotlib.axes._subplots.AxesSubplot at 0x26e0cc0bb70>
```



The performance of the model on the test set is highly satisfactory, as it demonstrates superior results with no instances of fraud being missed.